

Multithreaded Programming in

Cilk

LECTURE 2

Charles E. Leiserson

Supercomputing Technologies Research Group

Computer Science and Artificial Intelligence Laboratory
Massachusetts Institute of Technology

Minicourse Outline

- LECTURE 1

Basic Cilk programming: Cilk keywords, performance measures, scheduling.

- LECTURE 2

Analysis of Cilk algorithms: matrix multiplication, sorting, tableau construction.

- LABORATORY

Programming matrix multiplication in Cilk
— *Dr. Bradley C. Kuszmaul*

- LECTURE 3

Advanced Cilk programming: inlets, abort, speculation, data synchronization, & more.

LECTURE 2

- **Recurrences (Review)**
- **Matrix Multiplication**
- **Merge Sort**
- **Tableau Construction**
- **Conclusion**

The Master Method

The *Master Method* for solving recurrences applies to recurrences of the form

$$T(n) = aT(n/b) + f(n),^*$$

where $a \geq 1$, $b > 1$, and f is asymptotically positive.

IDEA: Compare $n^{\log_b a}$ with $f(n)$.

*The unstated base case is $T(n) = \Theta(1)$ for sufficiently small n .

Master Method — CASE 1

$$T(n) = aT(n/b) + f(n)$$

$$n^{\log_b a} \gg f(n)$$

Specifically, $f(n) = O(n^{\log_b a - \varepsilon})$ for some constant $\varepsilon > 0$.

Solution: $T(n) = \Theta(n^{\log_b a})$.

Master Method — CASE 2

$$T(n) = aT(n/b) + f(n)$$

$$n^{\log_b a} \approx f(n)$$

Specifically, $f(n) = \Theta(n^{\log_b a} \lg^k n)$ for some constant $k \geq 0$.

Solution: $T(n) = \Theta(n^{\log_b a} \lg^{k+1} n)$.

Master Method — CASE 3

$$T(n) = aT(n/b) + f(n)$$

$$n^{\log_b a} \ll f(n)$$

Specifically, $f(n) = \Omega(n^{\log_b a + \varepsilon})$ for some constant $\varepsilon > 0$ and $f(n)$ satisfies the *regularity condition* that $af(n/b) \leq cf(n)$ for some constant $c < 1$.

Solution: $T(n) = \Theta(f(n))$.

Master Method Summary

$$T(n) = a T(n/b) + f(n)$$

CASE 1: $f(n) = O(n^{\log_b a - \varepsilon})$, constant $\varepsilon > 0$

$$\Rightarrow T(n) = \Theta(n^{\log_b a}) .$$

CASE 2: $f(n) = \Theta(n^{\log_b a} \lg^k n)$, constant $k \geq 0$

$$\Rightarrow T(n) = \Theta(n^{\log_b a} \lg^{k+1} n) .$$

CASE 3: $f(n) = \Omega(n^{\log_b a + \varepsilon})$, constant $\varepsilon > 0$,

and regularity condition

$$\Rightarrow T(n) = \Theta(f(n)) .$$

Master Method Quiz

- $T(n) = 4 T(n/2) + n$

$$n^{\log_b a} = n^2 \gg n \Rightarrow \text{CASE 1: } T(n) = \Theta(n^2).$$

- $T(n) = 4 T(n/2) + n^2$

$$n^{\log_b a} = n^2 = n^2 \lg^0 n \Rightarrow \text{CASE 2: } T(n) = \Theta(n^2 \lg n).$$

- $T(n) = 4 T(n/2) + n^3$

$$n^{\log_b a} = n^2 \ll n^3 \Rightarrow \text{CASE 3: } T(n) = \Theta(n^3).$$

- $T(n) = 4 T(n/2) + n^2 / \lg n$

Master method does not apply!

LECTURE 2

- Recurrences (Review)
- **Matrix Multiplication**
- **Merge Sort**
- **Tableau Construction**
- **Conclusion**

Square-Matrix Multiplication

$$\begin{matrix} \begin{pmatrix} c_{11} & c_{12} & \cdots & c_{1n} \\ c_{21} & c_{22} & \cdots & c_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ c_{n1} & c_{n2} & \cdots & c_{nn} \end{pmatrix} & = & \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{pmatrix} & \times & \begin{pmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{nn} \end{pmatrix} \\ C & & A & & B \end{matrix}$$

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

Assume for simplicity that $n = 2^k$.

Recursive Matrix Multiplication

Divide and conquer —

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$
$$= \begin{bmatrix} A_{11}B_{11} & A_{11}B_{12} \\ A_{21}B_{11} & A_{21}B_{12} \end{bmatrix} + \begin{bmatrix} A_{12}B_{21} & A_{12}B_{22} \\ A_{22}B_{21} & A_{22}B_{22} \end{bmatrix}$$

8 multiplications of $(n/2) \times (n/2)$ matrices.

1 addition of $n \times n$ matrices.

Matrix Multiply in Pseudo-Cilk

```
cilk void Mult(*C, *A, *B, n) {  
    float *T = Cilk_alloc(n*n*sizeof(float));  
    < base case & partition matrices >  
    spawn Mult(C11,A11,B11,n/2);  
    spawn Mult(C12,A11,B12,n/2);  
    spawn Mult(C22,A21,B12,n/2);  
    spawn Mult(C21,A21,B11,n/2);  
    spawn Mult(T11,A12,B21,n/2);  
    spawn Mult(T12,A12,B22,n/2);  
    spawn Mult(T22,A22,B22,n/2);  
    spawn Mult(T21,A22,B21,n/2);  
    sync;  
    spawn Add(C,T,n);  
    sync;  
    return;  
}
```

$$C = A \cdot B$$

*Absence of type
declarations.*

Matrix Multiply in Pseudo-Cilk

```
cilk void Mult(*C, *A, *B, n) {  
    float *T = Cilk_alloc(n*n*sizeof(float));  
    < base case & partition matrices >  
    spawn Mult(C11,A11,B11,n/2);  
    spawn Mult(C12,A11,B12,n/2);  
    spawn Mult(C21,A21,B11,n/2);  
    spawn Mult(C22,A21,B12,n/2);  
    spawn Mult(T11,A12,B21,n/2);  
    spawn Mult(T12,A12,B22,n/2);  
    spawn Mult(T21,A22,B21,n/2);  
    spawn Mult(T22,A22,B22,n/2);  
    sync;  
    spawn Add(C,T,n);  
    sync;  
    return;  
}
```

$$C = A \cdot B$$

*Coarsen base cases
for efficiency.*

Matrix Multiply in Pseudo-Cilk

```
cilk void Mult(*C, *A, *B, n) {  
    float *T = Cilk_alloc(n*n*sizeof(float));  
    < base case & partition matrices >  
    spawn Mult(C11,A11,B11,n/2);  
    spawn Mult(C12,A11,B12,n/2);  
    spawn Mult(C22,A21,B12,n/2);  
    spawn Mult(C21,A21,B21,n/2);  
    spawn Mult(T11,A12,B21,n/2);  
    spawn Mult(T12,A12,B22,n/2);  
    spawn Mult(T22,A22,B22,n/2);  
    spawn Mult(T21,A22,B21,n/2);  
    sync;  
    spawn Add(C,T,n);  
    sync;  
    return;  
}
```

Also need a row-size argument for array indexing.

Submatrices are produced by pointer calculation, not copying of elements.

$$C = A \cdot B$$

Matrix Multiply in Pseudo-Cilk

```
cilk void Mult(*C, *A, *B, n) {  
    float *T = Cilk_alloc(n*n*sizeof(float));  
    < base case & partition matrices >  
    spawn Mult(C11,A11,B11,n/2);  
    spawn Mult(C12,A11,B12,n/2);  
    spawn Mult(C22,A21,B12,n/2);  
    spawn Mult(C21,A21,B11,n/2);  
    spawn Mult(T11,A12,B21,n/2);  
    spawn Mult(T12,A12,B22,n/2);  
    spawn Mult(T22,A22,B22,n/2);  
    spawn Mult(T21,A22,B21,n/2);  
    sync;  
    spawn Add(C,T,n);  
    sync;  
    return;  
}
```

$$C = A \cdot B$$

$$C = C + T$$

```
cilk void Add(*C, *T, n) {  
    < base case & partition matrices >  
    spawn Add(C11,T11,n/2);  
    spawn Add(C12,T12,n/2);  
    spawn Add(C21,T21,n/2);  
    spawn Add(C22,T22,n/2);  
    sync;  
    return;  
}
```


Work of Matrix Addition

```
cilk void Add(*C, *T, n) {  
    < base case & partition matrices >  
    spawn Add(C11, T11, n/2);  
    spawn Add(C12, T12, n/2);  
    spawn Add(C21, T21, n/2);  
    spawn Add(C22, T22, n/2);  
    sync;  
    return;  
}
```

Work: $A_1(n) = 4A_1(n/2) + \Theta(1)$
 $= \Theta(n^2)$ — **CASE 1**

$$n^{\log_b a} = n^{\log_2 4} = n^2 \gg \Theta(1).$$

Span of Matrix Addition

maximum

```
cilk void Add(*C, *T, n) {  
    < base case & partition matrices >  
    spawn Add(C11, T11, n/2);  
    spawn Add(C12, T12, n/2);  
    spawn Add(C21, T21, n/2);  
    spawn Add(C22, T22, n/2);  
    sync;  
    return;  
}
```

$$\begin{aligned} \textit{Span: } A_{\infty}(n) &= A_{\infty}(n/2) + \Theta(1) \\ &= \Theta(\lg n) \text{ --- } \mathbf{CASE\ 2} \end{aligned}$$

$$n^{\log_b a} = n^{\log_2 1} = 1 \Rightarrow f(n) = \Theta(n^{\log_b a} \lg^0 n) .$$

Work of Matrix Multiplication

8 {

```
cilk void Mult(*C, *A, *B, n) {  
    float *T = Cilk_alloca(n*n*sizeof(float));  
    < base case & partition matrices >  
    spawn Mult(C11,A11,B11,n/2);  
    spawn Mult(C12,A11,B12,n/2);  
    ⋮  
    spawn Mult(T21,A22,B21,n/2);  
    sync;  
    spawn Add(C,T,n);  
    sync;  
    return;  
}
```

$$\begin{aligned} \text{Work: } M_1(n) &= 8M_1(n/2) + A_1(n) + \Theta(1) \\ &= 8M_1(n/2) + \Theta(n^2) \\ &= \Theta(n^3) \quad \text{--- CASE 1} \end{aligned}$$

$$n^{\log_b a} = n^{\log_2 8} = n^3 \gg \Theta(n^2).$$

Span of Matrix Multiplication

8

```
cilk void Mult(*C, *A, *B, n) {  
    float *T = Cilk_alloc(n*n*sizeof(float));  
    < base case & partition matrices >  
    {  
        spawn Mult(C11,A11,B11,n/2);  
        spawn Mult(C12,A11,B12,n/2);  
        :  
        spawn Mult(T21,A22,B21,n/2);  
        sync;  
        spawn Add(C,T,n);  
        sync;  
        return;  
    }  
}
```

$$\begin{aligned} \text{Span: } M_{\infty}(n) &= M_{\infty}(n/2) + A_{\infty}(n) + \Theta(1) \\ &= M_{\infty}(n/2) + \Theta(\lg n) \\ &= \Theta(\lg^2 n) \quad \text{--- CASE 2} \end{aligned}$$

$$n^{\log_b a} = n^{\log_2 1} = 1 \Rightarrow f(n) = \Theta(n^{\log_b a} \lg^1 n) .$$

Parallelism of Matrix Multiply

Work: $M_1(n) = \Theta(n^3)$

Span: $M_\infty(n) = \Theta(\lg^2 n)$

Parallelism: $\frac{M_1(n)}{M_\infty(n)} = \Theta(n^3 / \lg^2 n)$

For 1000×1000 matrices,
parallelism $\approx (10^3)^3 / 10^2 = 10^7$.

Stack Temporaries

```
cilk void Mult(*C, *A, *B, n) {  
    float *T = Cilk_alloca(n*n*sizeof(float));  
    < base case & partition matrices >  
    spawn Mult(C11,A11,B11,n/2);  
    spawn Mult(C12,A11,B12,n/2);  
    ⋮  
    spawn Mult(T21,A22,B21,n/2);  
    sync;  
    spawn Add(C,T,n);  
    sync;  
    return;  
}
```

In hierarchical-memory machines (especially chip multiprocessors), memory accesses are so expensive that minimizing storage often yields higher performance.

IDEA: Trade off parallelism for less storage.

No-Temp Matrix Multiplication

```
cilk void MultA(*C, *A, *B, n) {  
    // C = C + A * B  
    < base case & partition matrices >  
    spawn MultA(C11,A11,B11,n/2);  
    spawn MultA(C12,A11,B12,n/2);  
    spawn MultA(C22,A21,B12,n/2);  
    spawn MultA(C21,A21,B11,n/2);  
    sync;  
    spawn MultA(C21,A22,B21,n/2);  
    spawn MultA(C22,A22,B22,n/2);  
    spawn MultA(C12,A12,B22,n/2);  
    spawn MultA(C11,A12,B21,n/2);  
    sync;  
    return;  
}
```

Saves space, but at what expense?

Work of No-Temp Multiply

```
cilk void MultA(*C, *A, *B, n) {  
    // C = C + A * B  
    < base case & partition matrices >  
    spawn MultA(C11,A11,B11,n/2);  
    spawn MultA(C12,A11,B12,n/2);  
    spawn MultA(C22,A21,B12,n/2);  
    spawn MultA(C21,A21,B11,n/2);  
    sync;  
    spawn MultA(C21,A22,B21,n/2);  
    spawn MultA(C22,A22,B22,n/2);  
    spawn MultA(C12,A12,B22,n/2);  
    spawn MultA(C11,A12,B21,n/2);  
    sync;  
    return;  
}
```

$$\begin{aligned} \text{Work: } M_1(n) &= 8 M_1(n/2) + \Theta(1) \\ &= \Theta(n^3) \quad \text{--- CASE 1} \end{aligned}$$

Span of No-Temp Multiply

maximum

```
cilk void MultA(*C, *A, *B, n) {  
    // C = C + A * B  
    < base case & partition matrices >  
    {  
        spawn MultA(C11,A11,B11,n/2);  
        spawn MultA(C12,A11,B12,n/2);  
        spawn MultA(C22,A21,B12,n/2);  
        spawn MultA(C21,A21,B11,n/2);  
        sync;  
        {  
            spawn MultA(C21,A22,B21,n/2);  
            spawn MultA(C22,A22,B22,n/2);  
            spawn MultA(C12,A12,B22,n/2);  
            spawn MultA(C11,A12,B21,n/2);  
            sync;  
            return;  
        }  
    }  
}
```

maximum

$$\begin{aligned} \text{Span: } M_{\infty}(n) &= 2 M_{\infty}(n/2) + \Theta(1) \\ &= \Theta(n) \quad \text{--- CASE 1} \end{aligned}$$

Parallelism of No-Temp Multiply

Work: $M_1(n) = \Theta(n^3)$

Span: $M_\infty(n) = \Theta(n)$

Parallelism: $\frac{M_1(n)}{M_\infty(n)} = \Theta(n^2)$

For 1000×1000 matrices,
parallelism $\approx (10^3)^3/10^3 = 10^6$.

Faster in practice!

Testing Synchronization

Cilk language feature: A programmer can check whether a Cilk procedure is “synched” (without actually performing a **sync**) by testing the pseudovariable **SYNCHED**:

- **SYNCHED** = 0 \Rightarrow some spawned children might not have returned.
- **SYNCHED** = 1 \Rightarrow all spawned children have definitely returned.

Best of Both Worlds

```
cilk void Mult1(*C, *A, *B, n) { // multiply & store
    { base case & partition matrices }
    spawn Mult1(C11,A11,B11,n/2); // multiply & store
    spawn Mult1(C12,A11,B12,n/2);
    spawn Mult1(C22,A21,B12,n/2);
    spawn Mult1(C21,A21,B11,n/2);
    if (SYNCHED) {
        spawn MultA1(C11,A12,B21,n/2); // multiply & add
        spawn MultA1(C12,A12,B22,n/2);
        spawn MultA1(C22,A22,B22,n/2);
        spawn MultA1(C21,A22,B21,n/2);
    } else {
        float *T = Cilk_alloca(n*n*sizeof(float));
        spawn Mult1(T11,A11,B11,n/2);
        spawn Mult1(T12,A11,B12,n/2);
        spawn Mult1(T21,A21,B11,n/2);
        spawn Mult1(T22,A21,B12,n/2);
        sync;
        spawn Add(C,T,T,n);
    }
    sync;
    return;
}
```

This code is just as parallel as the original, but it only uses more space if runtime parallelism actually exists.


Ordinary Matrix Multiplication

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

IDEA: Spawn n^2 inner products in parallel.
Compute each inner product in parallel.

Work: $\Theta(n^3)$
Span: $\Theta(\lg n)$
Parallelism: $\Theta(n^3/\lg n)$

BUT, this algorithm exhibits poor locality and does not exploit the cache hierarchy of modern microprocessors, especially CMP's.






LECTURE 2

- Recurrences (Review)
- Matrix Multiplication
- Merge Sort
- Tableau Construction
- Conclusion

Merging Two Sorted Arrays

```
void Merge(int *C, int *A, int *B, int na, int nb) {  
    while (na>0 && nb>0) {  
        if (*A <= *B) {  
            *C++ = *A++; na--;  
        } else {  
            *C++ = *B++; nb--;  
        }  
    }  
    while (na>0) {  
        *C++ = *A++; na--;  
    }  
    while (nb>0) {  
        *C++ = *B++; nb--;  
    }  
}
```

Time to merge n
elements = $\Theta(n)$.



Merge Sort

```
cilk void MergeSort(int *B, int *A, int n) {  
    if (n==1) {  
        B[0] = A[0];  
    } else {  
        int *C;  
        C = (int*) Cilk_alloc(n*sizeof(int));  
        spawn MergeSort(C, A, n/2);  
        spawn MergeSort(C+n/2, A+n/2, n-n/2);  
        sync;  
        Merge(B, C, C+n/2, n/2, n-n/2);  
    }  
}
```

merge

merge

merge

3 4 12 14 19 21 33 46

3 12 19 46

4 14 21 33

3 19

12 46

4 33

14 21

19 3 12 46 33 4 21 14

Work of Merge Sort

```
cilk void MergeSort(int *B, int *A, int n) {  
    if (n==1) {  
        B[0] = A[0];  
    } else {  
        int *C;  
        C = (int*) Cilk_alloc(n*sizeof(int));  
        spawn MergeSort(C, A, n/2);  
        spawn MergeSort(C+n/2, A+n/2, n-n/2);  
        sync;  
        Merge(B, C, C+n/2, n/2, n-n/2);  
    }  
}
```

$$\begin{aligned} \textit{Work: } T_1(n) &= 2T_1(n/2) + \Theta(n) \\ &= \Theta(n \lg n) \text{ — } \mathbf{CASE\ 2} \end{aligned}$$

$$n^{\log_b a} = n^{\log_2 2} = n \Rightarrow f(n) = \Theta(n^{\log_b a} \lg^0 n) .$$

Span of Merge Sort

```
cilk void MergeSort(int *B, int *A, int n) {  
    if (n==1) {  
        B[0] = A[0];  
    } else {  
        int *C;  
        C = (int*) Cilk_alloc(n*sizeof(int));  
        spawn MergeSort(C, A, n/2);  
        spawn MergeSort(C+n/2, A+n/2, n-n/2);  
        sync;  
        Merge(B, C, C+n/2, n/2, n-n/2);  
    }  
}
```

$$\begin{aligned} \textit{Span: } T_{\infty}(n) &= T_{\infty}(n/2) + \Theta(n) \\ &= \Theta(n) \quad \text{--- CASE 3} \end{aligned}$$

$$n^{\log_b a} = n^{\log_2 1} = 1 \ll \Theta(n) .$$

Parallelism of Merge Sort

Work: $T_1(n) = \Theta(n \lg n)$


Span: $T_\infty(n) = \Theta(n)$

PUNY!

Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta(\lg n)$

*We need to parallelize the **merge**!*

Parallel Merge



KEY IDEA: If the total number of elements to be merged in the two arrays is $n = na + nb$, the total number of elements in the larger of the two recursive merges is at most $(3/4)n$.

Parallel Merge

```
cilk void P_Merge(int *C, int *A, int *B,
                  int na, int nb) {
    if (na < nb) {
        spawn P_Merge(C, B, A, nb, na);
    } else if (na==1) {
        if (nb == 0) {
            C[0] = A[0];
        } else {
            C[0] = (A[0]<B[0]) ? A[0] : B[0]; /* minimum */
            C[1] = (A[0]<B[0]) ? B[0] : A[0]; /* maximum */
        }
    } else {
        int ma = na/2;
        int mb = BinarySearch(A[ma], B, nb);
        spawn P_Merge(C, A, B, ma, mb);
        spawn P_Merge(C+ma+mb, A+ma, B+mb, na-ma, nb-mb);
        sync;
    }
}
```

Coarsen base cases for efficiency.

Span of P_Merge

```
cilk void P_Merge(int *C, int *A, int *B,
                  int na, int nb) {
    if (na < nb) {
        .
        .
        .
    } else {
        int ma = na/2;
        int mb = BinarySearch(A[ma], B, nb);
        spawn P_Merge(C, A, B, ma, mb);
        spawn P_Merge(C+ma+mb, A+ma, B+mb, na-ma, nb-mb);
        sync;
    }
}
```

$$\begin{aligned} \text{Span: } T_{\infty}(n) &= T_{\infty}(3n/4) + \Theta(\lg n) \\ &= \Theta(\lg^2 n) \quad \text{--- CASE 2} \end{aligned}$$

$$n^{\log_b a} = n^{\log_{4/3} 1} = 1 \Rightarrow f(n) = \Theta(n^{\log_b a} \lg^1 n) .$$

Work of P_Merge

```
cilk void P_Merge(int *C, int *A, int *B,
                  int na, int nb) {
    if (na < nb) {
        .
        .
        .
    } else {
        int ma = na/2;
        int mb = BinarySearch(A[ma], B, nb);
        spawn P_Merge(C, A, B, ma, mb);
        spawn P_Merge(C+ma+mb, A+ma, B+mb, nb-mb);
        sync;
    }
}
```

HAIRY!

Work: $T_1(n) = T_1(\alpha n) + T_1((1-\alpha)n) + \Theta(\lg n)$,
where $1/4 \leq \alpha \leq 3/4$.

CLAIM: $T_1(n) = \Theta(n)$.

Analysis of Work Recurrence

$$T_1(n) = T_1(\alpha n) + T_1((1-\alpha)n) + \Theta(\lg n),$$

where $1/4 \leq \alpha \leq 3/4$.

Substitution method: Inductive hypothesis is $T_1(k) \leq c_1 k - c_2 \lg k$, where $c_1, c_2 > 0$. Prove that the relation holds, and solve for c_1 and c_2 .

$$\begin{aligned} T_1(n) &= T_1(\alpha n) + T_1((1-\alpha)n) + \Theta(\lg n) \\ &\leq c_1(\alpha n) - c_2 \lg(\alpha n) \\ &\quad + c_1((1-\alpha)n) - c_2 \lg((1-\alpha)n) + \Theta(\lg n) \end{aligned}$$

Analysis of Work Recurrence

$$T_1(n) = T_1(\alpha n) + T_1((1-\alpha)n) + \Theta(\lg n),$$

where $1/4 \leq \alpha \leq 3/4$.

$$\begin{aligned} T_1(n) &= T_1(\alpha n) + T_1((1-\alpha)n) + \Theta(\lg n) \\ &\leq c_1(\alpha n) - c_2 \lg(\alpha n) \\ &\quad + c_1(1-\alpha)n - c_2 \lg((1-\alpha)n) + \Theta(\lg n) \end{aligned}$$

Analysis of Work Recurrence

$$T_1(n) = T_1(\alpha n) + T_1((1-\alpha)n) + \Theta(\lg n),$$

where $1/4 \leq \alpha \leq 3/4$.

$$\begin{aligned} T_1(n) &= T_1(\alpha n) + T_1((1-\alpha)n) + \Theta(\lg n) \\ &\leq c_1(\alpha n) - c_2 \lg(\alpha n) \\ &\quad + c_1(1-\alpha)n - c_2 \lg((1-\alpha)n) + \Theta(\lg n) \\ &\leq c_1 n - c_2 \lg(\alpha n) - c_2 \lg((1-\alpha)n) + \Theta(\lg n) \\ &\leq c_1 n - c_2 (\lg(\alpha(1-\alpha)) + 2 \lg n) + \Theta(\lg n) \\ &\leq c_1 n - c_2 \lg n \\ &\quad - (c_2(\lg n + \lg(\alpha(1-\alpha))) - \Theta(\lg n)) \\ &\leq c_1 n - c_2 \lg n \end{aligned}$$

by choosing c_1 and c_2 large enough.

Parallelism of **P_Merge**

Work: $T_1(n) = \Theta(n)$

Span: $T_\infty(n) = \Theta(\lg^2 n)$

Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta(n/\lg^2 n)$

Parallel Merge Sort

```
cilk void P_MergeSort(int *B, int *A, int n) {  
    if (n==1) {  
        B[0] = A[0];  
    } else {  
        int *C;  
        C = (int*) Cilk_alloca(n*sizeof(int));  
        spawn P_MergeSort(C, A, n/2);  
        spawn P_MergeSort(C+n/2, A+n/2, n-n/2);  
        sync;  
        spawn P_Merge(B, C, C+n/2, n/2, n-n/2);  
    }  
}
```

Work of Parallel Merge Sort

```
cilk void P_MergeSort(int *B, int *A, int n) {  
    if (n==1) {  
        B[0] = A[0];  
    } else {  
        int *C;  
        C = (int*) Cilk_alloc(n*sizeof(int));  
        spawn P_MergeSort(C, A, n/2);  
        spawn P_MergeSort(C+n/2, A+n/2, n-n/2);  
        sync;  
        spawn P_Merge(B, C, C+n/2, n/2, n-n/2);  
    }  
}
```

$$\begin{aligned} \textit{Work: } T_1(n) &= 2 T_1(n/2) + \Theta(n) \\ &= \Theta(n \lg n) \text{ — } \mathbf{CASE\ 2} \end{aligned}$$

Span of Parallel Merge Sort

```
cilk void P_MergeSort(int *B, int *A, int n) {  
    if (n==1) {  
        B[0] = A[0];  
    } else {  
        int *C;  
        C = (int*) Cilk_alloc(n*sizeof(int));  
        spawn P_MergeSort(C, A, n/2);  
        spawn P_MergeSort(C+n/2, A+n/2, n-n/2);  
        sync;  
        spawn P_Merge(B, C, C+n/2, n/2, n-n/2);  
    }  
}
```

$$\begin{aligned} \textit{Span: } T_{\infty}(n) &= T_{\infty}(n/2) + \Theta(\lg^2 n) \\ &= \Theta(\lg^3 n) \quad \text{--- CASE 2} \end{aligned}$$

$$n^{\log_b a} = n^{\log_2 1} = 1 \Rightarrow f(n) = \Theta(n^{\log_b a} \lg^2 n) .$$

Parallelism of Merge Sort

Work: $T_1(n) = \Theta(n \lg n)$

Span: $T_\infty(n) = \Theta(\lg^3 n)$

Parallelism: $\frac{T_1(n)}{T_\infty(n)} = \Theta(n/\lg^2 n)$

LECTURE 2

- Recurrences (Review)
- Matrix Multiplication
- Merge Sort
- **Tableau Construction**
- **Conclusion**

Tableau Construction

Problem: Fill in an $n \times n$ tableau A , where $A[i, j] = f(A[i, j-1], A[i-1, j], A[i-1, j-1])$.


00	01	02	03	04	05	06	07
10	11	12	13	14	15	16	17
20	21	22	23	24	25	26	27
30	31	32	33	34	35	36	37
40	41	42	43	44	45	46	47
50	51	52	53	54	55	56	57
60	61	62	63	64	65	66	67
70	71	72	73	74	75	76	77

Dynamic programming

- Longest common subsequence
- Edit distance
- Time warping

Work: $\Theta(n^2)$.


Recursive Construction



Cilk code

```
spawn I;  
sync;  
spawn II;  
spawn III;  
sync;  
spawn IV;  
sync;
```

Recursive Construction




Cilk code

```
spawn I;  
sync;  
spawn II;  
spawn III;  
sync;  
spawn IV;  
sync;
```

$$\begin{aligned} \text{Work: } T_1(n) &= 4T_1(n/2) + \Theta(1) \\ &= \Theta(n^2) \quad \text{--- CASE 1} \end{aligned}$$

Recursive Construction



Cilk code

```
spawn I;  
sync;  
spawn II;  
spawn III;  
sync;  
spawn IV;  
sync;
```

$$\begin{aligned} \text{Span: } T_{\infty}(n) &= 3T_{\infty}(n/2) + \Theta(1) \\ &= \Theta(n^{\lg 3}) \text{ — CASE 1} \end{aligned}$$


Analysis of Tableau Construction

$$\textit{Work: } T_1(n) = \Theta(n^2)$$

$$\begin{aligned}\textit{Span: } T_\infty(n) &= \Theta(n^{\lg 3}) \\ &\approx \Theta(n^{1.58})\end{aligned}$$


$$\textit{Parallelism: } \frac{T_1(n)}{T_\infty(n)} \approx \Theta(n^{0.42})$$

A More-Parallel Construction



```
spawn I;  
sync;  
spawn II;  
spawn III;  
sync;  
spawn IV;  
spawn V;  
spawn VI;  
sync;  
spawn VII;  
spawn VIII;  
sync;  
spawn IX;  
sync;
```

A More-Parallel Construction




```

spawn I;
sync;
spawn II;
spawn III;
sync;
spawn IV;
spawn V;
spawn VI;
sync;
spawn VII;
spawn VIII;
sync;
spawn IX;
sync;
    
```

$$\begin{aligned}
 \text{Work: } T_1(n) &= 9T_1(n/3) + \Theta(1) \\
 &= \Theta(n^2) \quad \text{--- CASE 1}
 \end{aligned}$$

A More-Parallel Construction



```

spawn I;
sync;
spawn II;
spawn III;
sync;
spawn IV;
spawn V;
spawn VI;
sync;
spawn VII;
spawn VIII;
sync;
spawn IX;
sync;
    
```

$$\begin{aligned}
 \text{Span: } T_{\infty}(n) &= 5T_{\infty}(n/3) + \Theta(1) \\
 &= \Theta(n^{\log_3 5}) \text{ — CASE 1}
 \end{aligned}$$

Analysis of Revised Construction

$$\textit{Work: } T_1(n) = \Theta(n^2)$$

$$\begin{aligned}\textit{Span: } T_\infty(n) &= \Theta(n^{\log_3 5}) \\ &\approx \Theta(n^{1.46})\end{aligned}$$

$$\textit{Parallelism: } \frac{T_1(n)}{T_\infty(n)} \approx \Theta(n^{0.54})$$

More parallel by a factor of

$$\Theta(n^{0.54}) / \Theta(n^{0.42}) = \Theta(n^{0.12}) .$$

Puzzle

What is the largest parallelism that can be obtained for the tableau-construction problem using Cilk?

- You may only use basic Cilk control constructs (**spawn**, **sync**) for synchronization.
- No locks, synchronizing through memory, etc.

LECTURE 2

- Recurrences (Review)
- Matrix Multiplication
- Merge Sort
- Tableau Construction
- **Conclusion**

Key Ideas

- Cilk is simple: **cilk**, **spawn**, **sync**, **SYNCHED**
- Recurrences, recurrences, recurrences, ...

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span


- Work & span

- Work & span

- Work & span

- Work & span

- Work & span



Minicourse Outline

- **LECTURE 1**

Basic Cilk programming: Cilk keywords, performance measures, scheduling.

- **LECTURE 2**

Analysis of Cilk algorithms: matrix multiplication, sorting, tableau construction.

- **LABORATORY**

Programming matrix multiplication in Cilk
— *Dr. Bradley C. Kuszmaul*

- **LECTURE 3**

Advanced Cilk programming: inlets, abort, speculation, data synchronization, & more.