

Multithreaded Programming in

Cilk

LECTURE 3

Charles E. Leiserson

Supercomputing Technologies Research Group

Computer Science and Artificial Intelligence Laboratory
Massachusetts Institute of Technology

Minicourse Outline

- **LECTURE 1**

Basic Cilk programming: Cilk keywords, performance measures, scheduling.

- **LECTURE 2**

Analysis of Cilk algorithms: matrix multiplication, sorting, tableau construction.

- **LABORATORY**

Programming matrix multiplication in Cilk
— *Dr. Bradley C. Kuszmaul*

- **LECTURE 3**

Advanced Cilk programming: inlets, abort, speculation, data synchronization, & more.

LECTURE 3

- **Inlets**
- **Abort**
- **Speculative Computing**
- **Data Synchronization**
- **Under the Covers**
- **JCilk**
- **Conclusion**

Operating on Returned Values

Programmers may sometimes wish to incorporate a value returned from a spawned child into the parent frame by means other than a simple variable assignment.

Example:

```
x += spawn foo(a,b,c);
```

Cilk achieves this functionality using an internal function, called an *inlet*, which is executed as a secondary thread on the parent frame when the child returns.

Semantics of Inlets

```
int max, ix = -1;
inlet void update ( int val, int index ) {
    if (idx == -1 || val > max) {
        ix = index; max = val;
    }
}
...
for (i=0; i<1000000; i++) {
    update ( spawn foo(i), i );
}
sync; /* ix now indexes the largest foo(i) */
```

- The **inlet** keyword defines a **void** internal function to be an inlet.
- In the current implementation of Cilk, the inlet definition may not contain a **spawn**, and only the first argument of the inlet may be spawned at the call site.

Semantics of Inlets

```
int max, ix = -1;
inlet void update ( int val, int index ) {
    if (idx == -1 || val > max) {
        ix = index; max = val;
    }
}
...
for (i=0; i<1000000; i++) {
    update ( spawn foo(i), i );
}
sync; /* ix now indexes the largest foo(i) */
```

1. The non-**spawn** args to **update()** are evaluated.
2. The Cilk procedure **foo(i)** is spawned.
3. Control passes to the next statement.
4. When **foo(i)** returns, **update()** is invoked.

Semantics of Inlets

```
int max, ix = -1;
inlet void update ( int val, int index ) {
    if (idx == -1 || val > max) {
        ix = index; max = val;
    }
}
...
for (i=0; i<1000000; i++) {
    update ( spawn foo(i), i );
}
sync; /* ix now indexes the largest foo(i) */
```

Cilk provides *implicit atomicity* among the threads belonging to the same frame, and thus no locking is necessary to avoid data races.

Implicit Inlets

```
cilk int wfib(int n) {  
    if (n == 0) {  
        return 0;  
    } else {  
        int i, x = 1;  
        for (i=0; i<=n-2; i++) {  
            x += spawn wfib(i);  
        }  
        sync;  
        return x;  
    }  
}
```

For assignment operators, the Cilk compiler automatically generates an *implicit inlet* to perform the update.

LECTURE 3

- Inlets
- **Abort**
- **Speculative Computing**
- **Data Synchronization**
- **Under the Covers**
- **JCilk**
- **Conclusion**

Computing a Product

$$p = \prod_{i=0}^n A_i$$

```
int product(int *A, int n) {  
    int i, p=1;  
    for (i=0; i<n; i++) {  
        p *= A[i];  
    }  
    return p;  
}
```

Optimization: Quit early if the partial product ever becomes 0.

Computing a Product

$$p = \prod_{i=0}^n A_i$$

```
int product(int *A, int n) {  
    int i, p=1;  
    for (i=0; i<n; i++) {  
        p *= A[i];  
        if (p == 0) break;  
    }  
    return p;  
}
```

Optimization: Quit early if the partial product ever becomes 0.

Computing a Product in Parallel

$$p = \prod_{i=0}^n A_i$$

```
cilk int prod(int *A, int n) {  
    int p = 1;  
    if (n == 1) {  
        return A[0];  
    } else {  
        p *= spawn product(A, n/2);  
        p *= spawn product(A+n/2, n-n/2);  
        sync;  
        return p;  
    }  
}
```

How do we quit early if we discover a zero?

Cilk's Abort Feature

```
cilk int product(int *A, int n) {
    int p = 1;
    inlet void mult(int x) {
        p *= x;
        return;
    }

    if (n == 1) {
        return A[0];
    } else {
        mult( spawn product(A, n/2) );
        mult( spawn product(A+n/2, n-n/2) );
        sync;
        return p;
    }
}
```

1. Recode the implicit inlet to make it explicit.

Cilk's Abort Feature

```
cilk int product(int *A, int n) {  
    int p = 1;  
    inlet void mult(int x) {  
        p *= x;  
  
        return;  
    }  
  
    if (n == 1) {  
        return A[0];  
    } else {  
        mult( spawn product(A, n/2) );  
        mult( spawn product(A+n/2, n-n/2) );  
        sync;  
        return p;  
    }  
}
```

2. Check for 0 within the inlet.

Cilk's Abort Feature

```
cilk int product(int *A, int n) {
    int p = 1;
    inlet void mult(int x) {
        p *= x;
        if (p == 0) {
            abort; /* Aborts existing children, */
        }        /* but not future ones. */
        return;
    }

    if (n == 1) {
        return A[0];
    } else {
        mult( spawn product(A, n/2) );
        mult( spawn product(A+n/2, n-n/2) );
        sync;
        return p;
    }
}
```

2. Check for 0 within the inlet.

Cilk's Abort Feature

```
cilk int product(int *A, int n) {
    int p = 1;
    inlet void mult(int x) {
        p *= x;
        if (p == 0) {
            abort; /* Aborts existing children, */
        }        /* but not future ones. */
        return;
    }

    if (n == 1) {
        return A[0];
    } else {
        mult( spawn product(A, n/2) );

        mult( spawn product(A+n/2, n-n/2) );
        sync;
        return p;
    }
}
```


Cilk's Abort Feature

```
cilk int product(int *A, int n) {
    int p = 1;
    inlet void mult(int x) {
        p *= x;
        if (p == 0) {
            abort; /* Aborts existing children, */
                /* but not future ones. */
        }
        return;
    }


    if (n == 1) {
        return A[0];
    } else {
        mult( spawn product(A, n/2) );
        if (p == 0) { /* Don't spawn if we've */
            return 0; /* already aborted! */
        }
        mult( spawn product(A+n/2, n-n/2) );
        sync;
        return p;
    }
}
```

Implicit atomicity eases reasoning about races.

LECTURE 3


- Inlets
- Abort
- **Speculative Computing**
- **Data Synchronization**
- **Under the Covers**
- **JCilk**
- **Conclusion**

Min-Max Search




- Two players: MAX ■ and MIN ●.
- The game tree represents all moves from the current position within a given search depth.
- At leaves, apply a static evaluation function.
- MAX chooses the maximum score among its children.
- MIN chooses the minimum score among its children.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning



IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Alpha-Beta Pruning




IDEA: If MAX ■ discovers a move so good that MIN ● would never allow that position, MAX's other children need not be searched — *beta cutoff*.

Unfortunately, this heuristic is inherently serial.

Parallel Min-Max Search

OBSERVATION: In a best-ordered tree, the degree of every internal node is either **1** or *maximal*.



IDEA: [Feldman-Mysliwicz-Monien 91] If the first child fails to generate a cutoff, *speculate* that the remaining children can be searched in parallel without wasting any work: “*young brothers wait.*”

Parallel Alpha-Beta (I)

```
cilk int search(position *prev, int move, int depth) {  
    position cur;           /* Current position */  
    int bestscore = -INF;    /* Best score so far */  
    int num_moves;          /* Number of children */  
    int mv;                 /* Index of child */  
    int sc;                 /* Child's score */  
    int cutoff = FALSE;     /* Have we seen a cutoff? */
```

- View from MAX's perspective; MIN's viewpoint can be obtained by negating scores — *negamax*.
- The node generates its children from its parent's position **prev**.
- The **alpha** and **beta** limits and the move list are fields of the **position** data structure.

*#Cilk keywords
used so far*

1

Parallel Alpha-Beta (II)

```
inlet void get_score(int child_sc) {  
    child_sc = -child_sc;    /* Negamax */  
  
    if (child_sc > bestscore) {  
        bestscore = child_sc;  
        if (child_sc > cur.alpha) {  
            cur.alpha = child_sc;  
            if (child_sc >= cur.beta) { /* Beta cutoff */  
                cutoff = TRUE; /* No need to search more */  
                abort;         /* Terminate other children */  
            }  
        }  
    }  
}
```

Parallel Alpha-Beta (III)

```
/* Create current position and set up for search */  
  
make_move(prev, move, &cur);  
  
sc = eval(&cur);    /* Static evaluation */  
if ( abs(sc)>=MATE || depth<=0 ) { /* Leaf node */  
    return (sc);  
}  
  
cur.alpha = -prev->beta;    /* Negamax */  
cur.beta = -prev->alpha;  
  
/* Generate moves, hopefully in best-first order*/  
num_moves = gen_moves(&cur);
```

Parallel Alpha-Beta (IV)

```
/* Search the moves */

for (mv=0; !cutoff && mv<num_moves; mv++) {
    get_score( spawn search(&cur, mv, depth-1) );
    if (mv==0) sync; /* Young brothers wait */
}
sync;
return (bestscore);
}
```

- Only 6 Cilk keywords need be embedded in the C program to parallelize it.
- In fact, the program can be parallelized using only 5 keywords at the expense of minimal obfuscation.

LECTURE 3

- Inlets
- Abort
- Speculative Computing
- **Data Synchronization**
- **Under the Covers**
- **JCilk**
- **Conclusion**

Mutual Exclusion

Cilk's solution to mutual exclusion is no better than anybody else's.

Cilk provides a library of spin locks declared with **Cilk_lockvar**.

- To avoid deadlock with the Cilk scheduler, a lock should only be held within a Cilk thread.
- *I.e.*, **spawn** and **sync** should not be executed while a lock is held.

Fortunately, Cilk's control parallelism often mitigates the need for extensive locking.

Cilk's Memory Model


Programmers may also synchronize through memory using lock-free protocols, although Cilk is agnostic on consistency model.

- If a program contains no data races, Cilk effectively supports sequential consistency.
- If a program contains data races, Cilk's behavior depends on the consistency model of the underlying hardware.

To aid portability, the `Cilk_fence()` function implements a memory barrier on machines with weak memory models.

Debugging Data Races

Cilk's *Nondeterminator* debugging tool provably guarantees to detect and localize data-race bugs.

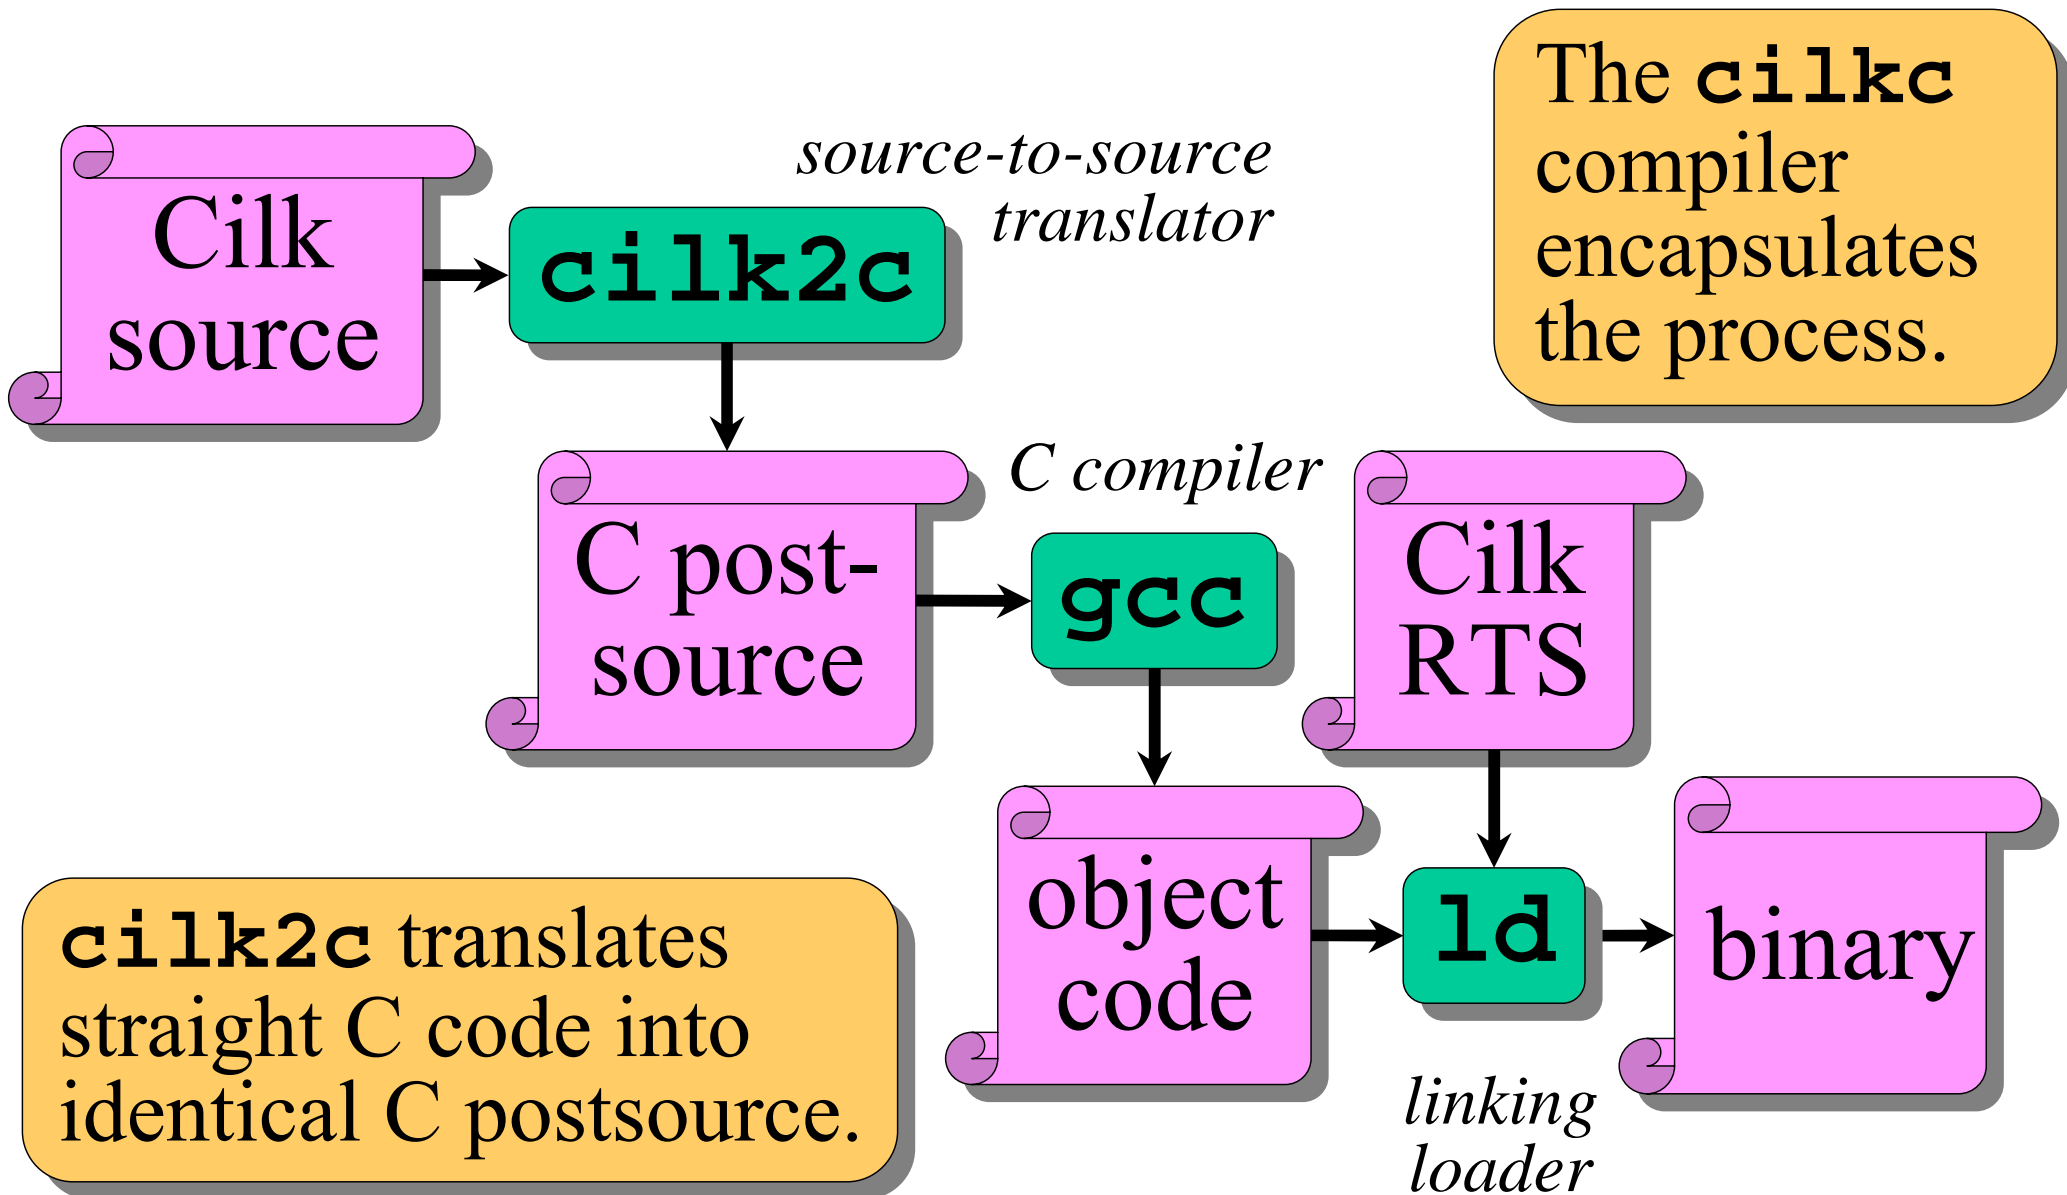


A *data race* occurs whenever two logically parallel threads, holding no locks in common, access the same location and one of the threads modifies the location.

LECTURE 3

- Inlets
- Abort
- Speculative Computing
- Data Synchronization
- Under the Covers
- JCilk
- Conclusion

Compiling Cilk




Cilk's Compiler Strategy

The **cilk2c** translator generates two “clones” of each Cilk procedure:

- *fast clone*—serial, common-case code.
- *slow clone*—code with parallel bookkeeping.

-
-
- The *fast clone* is always spawned, saving live variables on Cilk's work deque (shadow stack).
 - The *slow clone* is resumed if a thread is stolen, restoring variables from the shadow stack.
 - A check is made whenever a procedure returns to see if the resuming parent has been stolen.



Compiling **spawn** — Fast Clone

Cilk
source

```
x = spawn fib(n-1);
```




```
frame->entry = 1;  
frame->n = n;  
push(frame);
```

```
x = fib(n-1);
```

```
if (pop() == FAILURE) {  
    frame->x = x;  
    frame->join--;  
    { clean up &  
      return to scheduler }  
}
```

C post-
source

frame



} suspend
parent

} run child

} resume
parent
remotely

Cilk
deque

Compiling **sync** — Fast Clone

*Cilk
source*

sync ;

cilk2c

*C post-
source*

;

SLOW

FAST

FAST

FAST

FAST

FAST

No synchronization overhead in the fast clone!

Compiling the Slow Clone

```

void fib_slow(fib_frame *frame) {
    int n,x,y;
    switch (frame->entry) {
        case 1: goto L1;
        case 2: goto L2;
        case 3: goto L3;
    }
    :
    frame->entry = 1;
    frame->n = n;
    push(frame);
    x = fib(n-1);
    if (pop()==FAILURE) {
        frame->x = x;
        frame->join--;
        { clean up &
          return to scheduler }
    }

    if (0) {
        L1:;
        n = frame->n;
    }
    :
}

```


restore
program
counter

same
as fast
clone

restore local
variables
if resuming


continue

frame



*Cilk
deque*

Breakdown of Work Overhead (circa 1997)

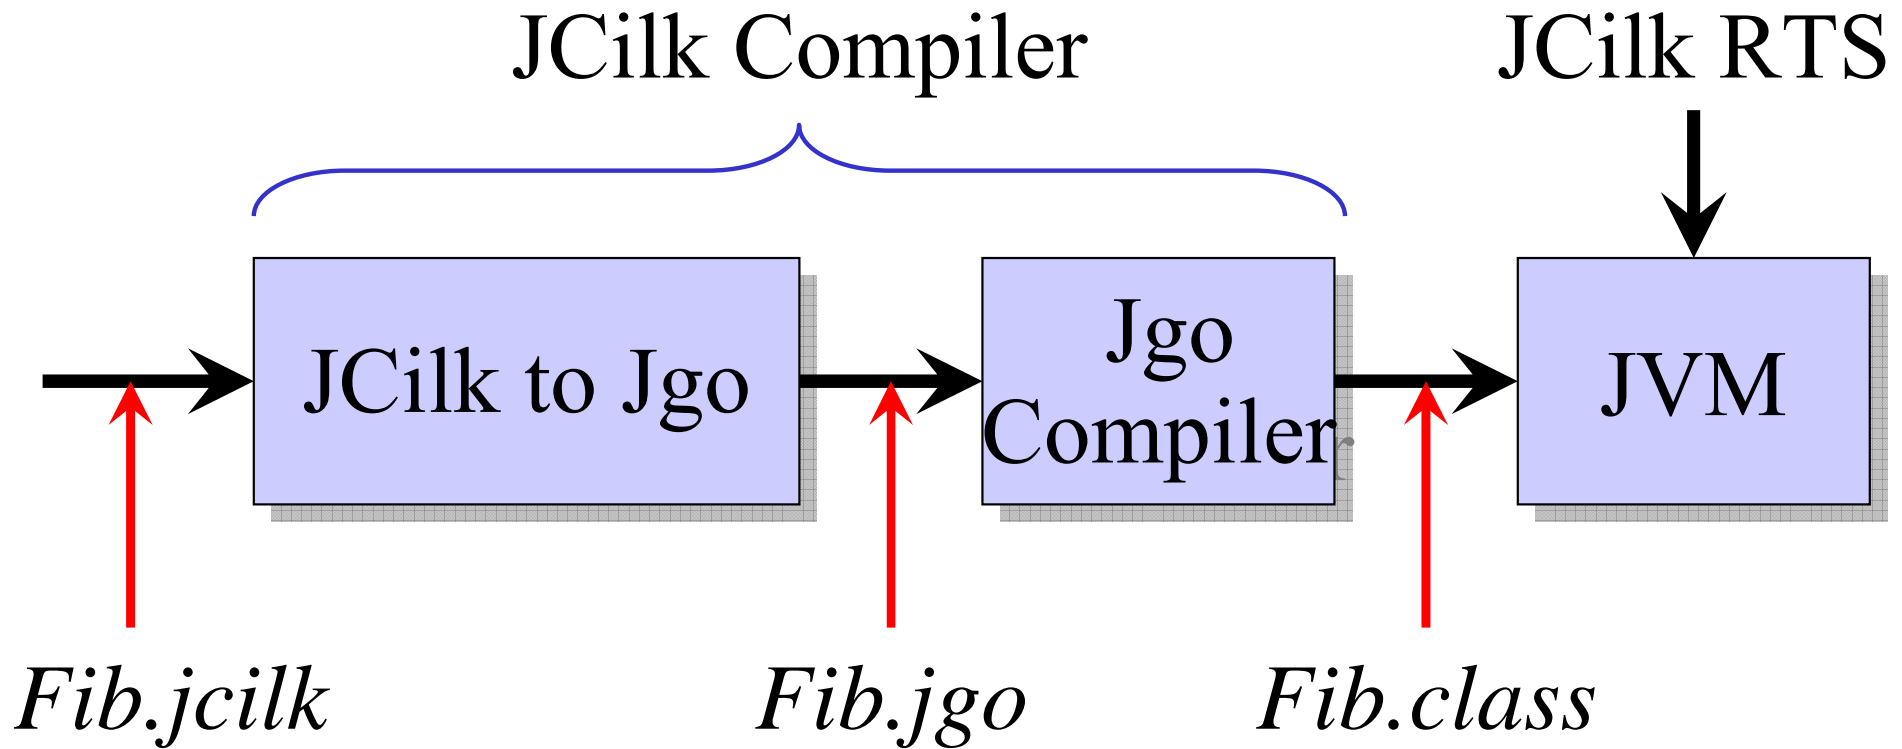


Benchmark: `fib` on one processor.

LECTURE 3

- Inlets
- Abort
- Speculative Computing
- Data Synchronization
- Under the Covers
- **JCilk**
- **Conclusion**

The JCilk System



- Jgo = Java + **goto**.
- The Jgo compiler was built by modifying **gcj** to accept **goto** statements so that a continuation mechanism for JCilk could be implemented.

JCilk Keywords

cilk

spawn

sync

SYNCHED

Same as Cilk, except that **cilk** can also modify **try**.

~~**inlet**~~

~~**abort**~~

Eliminated!

JCilk leverages Java's exception mechanism to render two Cilk keywords unnecessary.

Exception Handling in Java

“During the process of throwing an exception, the Java virtual machine *abruptly completes*, one by one, any expressions, statements, method and constructor invocations, initializers, and field initialization expressions that have begun but not completed execution in the current thread. This process continues until a handler is found that indicates that it handles that particular exception by naming the class of the exception or a superclass of the class of the exception.”

— J. Gosling, B Joy, G. Steele, and G. Bracha,
Java Language Specification, 2000, pp. 219–220.

Exception Handling in JCilk

```
private cilk void foo() throws IOException {  
    spawn A();  
    cilk try {  
        spawn B();  
        cilk try {  
            spawn C();  
        } catch(ArithmeticEx'n e) {  
            doSomething();  
        }  
    } catch(RuntimeException e) {  
        doSomethingElse();  
    }  
    spawn D();  
    doYetSomethingElse();  
    sync;  
}
```

Exception Handling in JCilk

```
private cilk void foo() throws IOException {  
    spawn A();  
    cilk try {  
        spawn B();  
        cilk try {  
            spawn C();  
        } catch(ArithmeticEx'n e) {  
            doSomething();  
        }  
    } catch(RuntimeException e) {  
        doSomethingElse();  
    }  
    spawn D();  
    doYetSomethingElse();  
    sync;  
}
```



Exception!

*An exception causes all
subcomputations
dynamically enclosed
by the catching clause
to abort!*

Exception Handling in JCilk

```
private cilk void foo() throws IOException {  
    spawn A();  
    cilk try {  
        spawn B();  
        cilk try {  
            spawn C();  
        } catch(ArithmeticEx'n e) {  
            doSomething();  
        }  
    } catch(RuntimeException e) {  
        doSomethingElse();  
    }  
    spawn D();  
    doYetSomethingElse();  
    sync;  
}
```



ArithmeticEx'n

Nothing
aborts.

*An exception causes all
subcomputations
dynamically enclosed
by the catching clause
to abort!*

Exception Handling in JCilk

```
private cilk void foo() throws IOException {  
    spawn A();  
    cilk try {  
        spawn B();  
        cilk try {  
            spawn C();  
        } catch(ArithmeticEx'n e) {  
            doSomething();  
        }  
    } catch(RuntimeException e) {  
        doSomethingElse();  
    }  
    spawn D();  
    doYetSomethingElse();  
    sync;  
}
```



RuntimeEx'n

*An exception causes all
subcomputations
dynamically enclosed
by the catching clause
to abort!*

Exception Handling in JCilk

```
private cilk void foo() throws IOException {  
    spawn A();  
    cilk try {  
        spawn B();  
        cilk try {  
            spawn C();  
        } catch(ArithmeticEx'n e) {  
            doSomething();  
        }  
    } catch(RuntimeException e) {  
        doSomethingElse();  
    }  
    spawn D();  
    doYetSomethingElse();  
    sync;  
}
```



IOException

An exception causes all subcomputations dynamically enclosed by the catching clause to abort!

Exception Handling in JCilk

```
private cilk void foo() throws IOException {  
    spawn A();  
    cilk try {  
        spawn B();  
        cilk try {  
            spawn C();  
        } catch(ArithmeticEx'n e) {  
            doSomething();  
        }  
    } catch(RuntimeException e) {  
        doSomethingElse();  
    }  
    spawn D();  
    doYetSomethingElse();  
    sync;  
}
```



RuntimeEx'n

*The appropriate **catch** clause is executed only after all spawned methods within the corresponding **try** block terminate.*

JCilk's Exception Mechanism

- JCilk's exception semantics allow programs such as alpha-beta to be coded without Cilk's **inlet** and **abort** keywords.
- Unfortunately, Java exceptions are slow, reducing the utility of JCilk's faithful extension.

LECTURE 3

- Inlets
- Abort
- Speculative Computing
- Data Synchronization
- Under the Covers
- JCilk
- **Conclusion**

Future Work

Adaptive computing

- Get rid of `--nproc` .
- Build a job scheduler that uses *parallelism feedback* to balance processor resources among Cilk jobs.

Integrating Cilk with static threads

- Currently, interfacing a Cilk program to other system processes requires arcane knowledge.
- Build linguistic support into Cilk for Cilk processes that communicate.
- Develop a job scheduler that uses *pipeload* to allocate resources among Cilk processes.

Key Ideas

- Cilk is simple: **cilk**, **spawn**, **sync**, **SYNCHED**, **inlet**, **abort**
- JCilk is simpler

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span

- Work & span


- Work & span

- Work & span

- Work & span

- Work & span

- Work & span



Open-Cilk Consortium

- We are in the process of forming a consortium to manage, organize, and promote Cilk open-source technology.
- If you are interested in participating, please let us know.

ACM Symposium on Parallelism in Algorithms and Architectures

SPAA 2006



Cambridge, MA, USA
July 30 – August 2, 2006