

CS341: Operating System**Process Scheduling**Lect12 : 27th Aug 2014

Dr. A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Outline

- Process Concepts Recap
 - Process States, PCB, Context Switch
 - **Job Queue, Ready Queue, De**
- Scheduling
 - **System Oriented and Theoretical Analysis**
 - **Introduction to Scheduling Algorithms**
 - **Real Time Scheduling Algorithms**
 - **Multiprocessor Scheduling Algorithms**
 - **Distributed and Power Aware Scheduling**

2

Scheduling Algorithm Optimization**Criteria**

- Maximize
 - CPU utilization
 - Throughput
- Minimize
 - Turnaround time
 - Waiting time
 - Response time

Assumption

- Task/Process/Job used interchangably
- Let all tasks in memory
- Let all tasks don't do any IPC
- All task are independent
- All tasks don't require any other resource other then CPU

Five Popular Scheduling

- First Come First Serve
- Shortest Job First
- Shortest Remaining Time First
 - SJF-I
- Round-Robin Scheduler
- Priority Scheduler
- Priority-I
- Multi-Level Priority Queue
- Feed Back Priority Queue

5

Problem Cases

- Blindness about job types
 - I/O goes idle
- Optimization involves favoring jobs of type "A" over "B".
 - Lots of A's? B's starve
- Interactive process trapped behind others.
 - Response time sucks for no reason
- Priority Inversion: A depends on B. A's priority > B's.
 - B never runs

FCFS Scheduling

| Process | Burst Time/ Execution Time |
|---------|-------------------------------|
| P1 | 24 |
| P2 | 3 |
| P3 | 3 |

- Assume: processes arrive in the order: P_1, P_2, P_3
The **Gantt's Chart/Gantt Chart** for schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time
- SJF is optimal** – gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
 - Could ask the user

Example of SJF

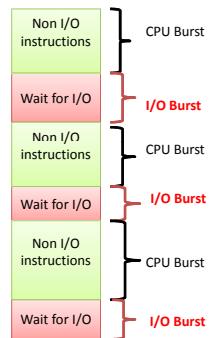
| Process | Burst Time |
|---------|------------|
| P1 | 6 |
| P2 | 8 |
| P3 | 7 |
| P4 | 3 |



$$\text{Average waiting time} = (3 + 16 + 9 + 0) / 4 = 7$$

Determining Length of Next CPU Burst

- Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using exponential averaging



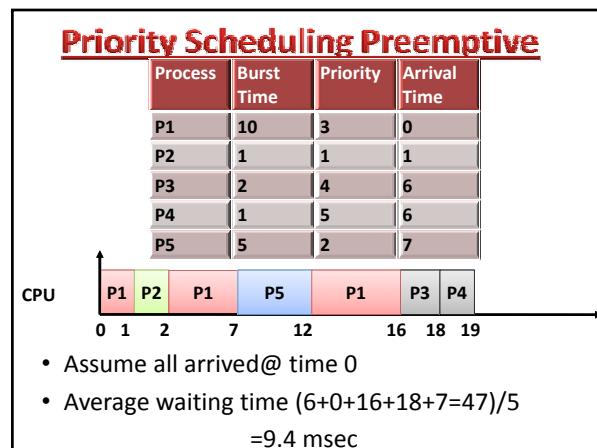
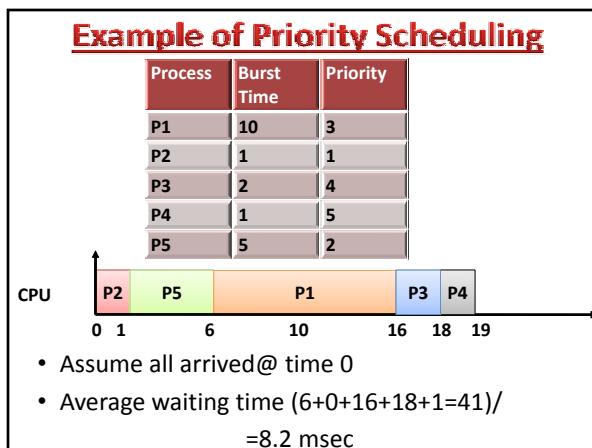
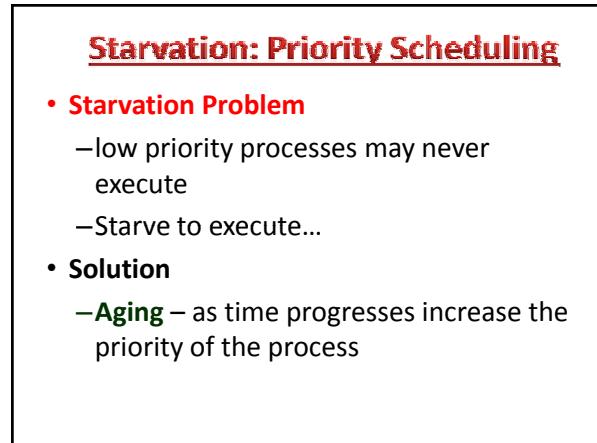
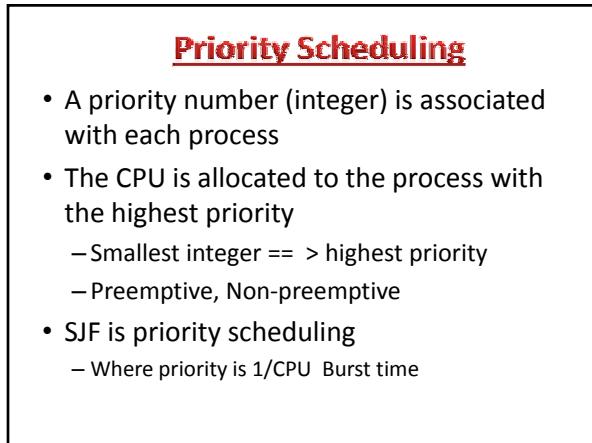
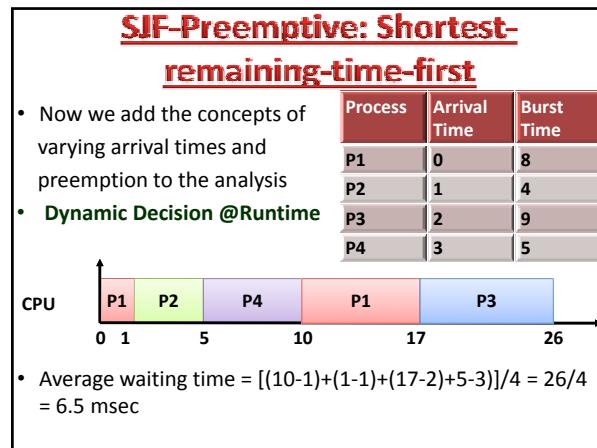
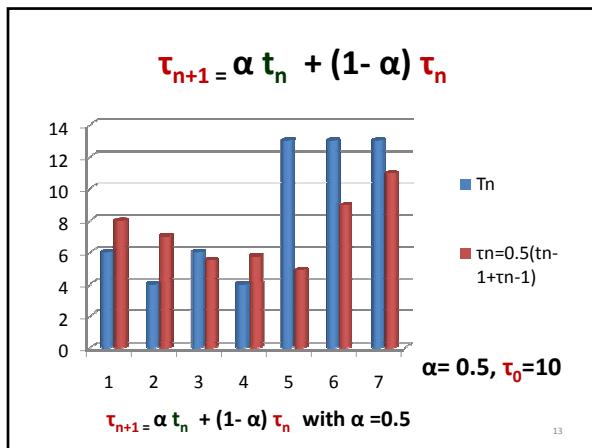
Prediction: Next CPU Burst

- Can be done by using the length of previous CPU bursts, using exponential averaging
 - t_n is actual CPU Burst of nth CPU Burst
 - τ_{n+1} = predicted values of next cpu burst
 - $-\alpha, 0 \leq \alpha \leq 1$
- $$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$
- Commonly, α set to $\frac{1}{2}$
- Preemptive version called **shortest-remaining-time-first**

Examples of Exponential Averaging

- $\alpha = 0 : \tau_{n+1} = \tau_n$, Recent history does not count
- $\alpha = 1: \tau_{n+1} = \alpha t_n$, Only actual last CPU burst counts
- If we expand the formula, we get:

$$\begin{aligned} \tau_{n+1} &= \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots \\ &\quad + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ &\quad + (1 - \alpha)^{n+1} \tau_0 \end{aligned}$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor



Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum q**)
 - Usually 10-100 milliseconds.
 - After this time has elapsed, the process is preempted and added to the end of the ready queue.

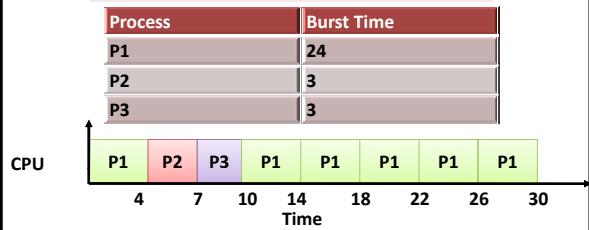
Round Robin (RR)

- If there are n processes in the ready queue and the time quantum is q , then
 - Each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
- No process waits more than $(n-1)q$ time units.

Round Robin (RR)

- Timer interrupts every quantum to schedule next process
 - Timer: Hardware unit, similar to Alarm
- Performance
 - q large \Rightarrow FIFO
 - q small \Rightarrow q must be large with respect to context switch, otherwise overhead is too high

Example : RR (with $q = 4$)



- Typically, higher Average turnaround than SJF, but better **response**
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec

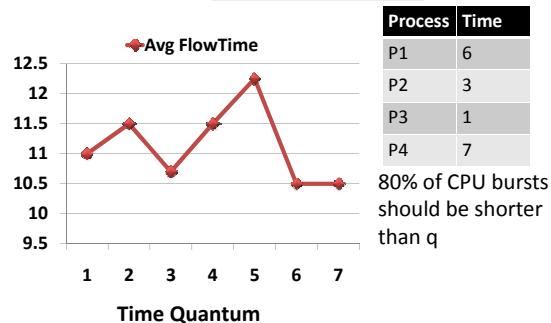
Time Quantum and Context Switch Time

Process time 10



| Quantum | Context Switches |
|---------|------------------|
| 12 | 0 |
| 6 | 1 |
| 1 | 9 |

Turnaround Time Varies With The Time Quantum



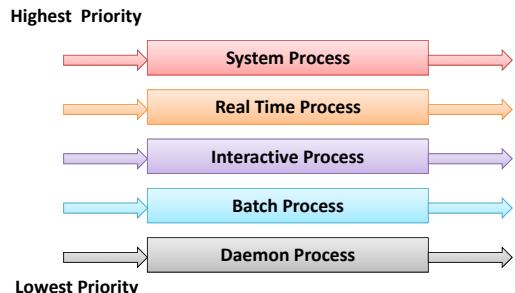
Multilevel Queue

- Ready queue is partitioned into separate queues, eg:
 - foreground** (interactive)
 - background** (batch)
- Process permanently in a given queue
- Each queue has its own scheduling algorithm:
 - foreground – RR
 - background – FCFS

Multilevel Queue

- Scheduling must be done between the queues:
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). **Possibility of starvation.**
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes;
 - i.e., 80% to foreground in RR, 20% to background in FCFS

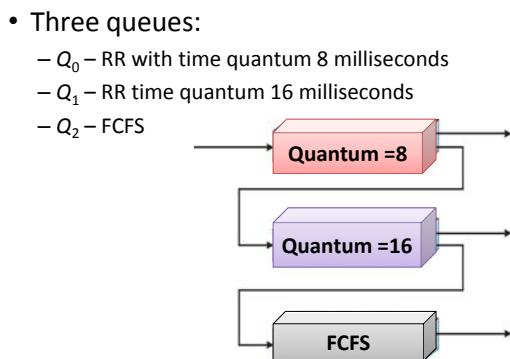
Multilevel Queue Scheduling



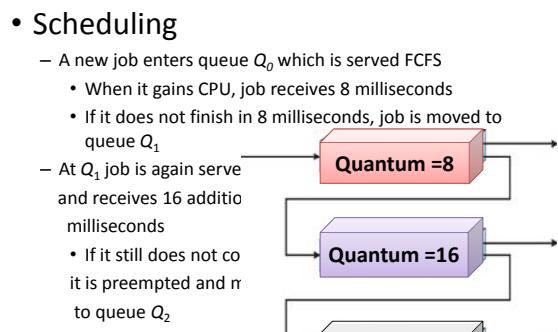
Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue



Example of Multilevel Feedback Queue



Introduction to Scheduling Algorithms

A bit of Theoretical View

31

Parallel Machine Problems

- For **identical machines** M_1, \dots, M_m
 - The processing time for j is the same on each machine.
- For **unrelated machines**
 - The processing time p_{jk} depends on the machine M_k on which j is processed.

Parallel Machine Problems

- For **uniform machine**
 - if $p_{jk} = p_j/r_k$.
- For **multi-purpose machines**
 - A set of machines μ_j is associated with each job j indicating that j can be processed on one machine in μ_j only.

Example: Machine Environment

| | M1 | M2 | M3 |
|----|----|----|----|
| P1 | 5 | 5 | 5 |
| P2 | 8 | 8 | 8 |
| P3 | 6 | 6 | 6 |

Identical

| | M1 | M2 | M3 |
|----|----|----|----|
| P1 | 5 | 4 | 6 |
| P2 | 9 | 8 | 4 |
| P3 | 3 | 18 | 4 |

Unrelated

| | M1 | M2 | M3 |
|----|----|-------|-----|
| P1 | 5 | 5/1.5 | 5/2 |
| P2 | 9 | 9/1.5 | 9/2 |
| P3 | 9 | 9/1.5 | 9/2 |

Uniform

34

Classification of Scheduling Problems

Classes of scheduling problems can be specified in terms of the three-field classification

$\alpha | \beta | \gamma$

where

- α specifies the **machine environment**,
- β specifies the **job characteristics**, and
- γ describes the **objective function(s)**.

Machine Environment

- **1** single machine
- **P** parallel identical machines
- **Q** uniform machines
- **R** unrelated machines
- MPM multipurpose machines, J job-shop,
- F flow-shop O open-shop

If the number of machines is fixed to m we write **Pm**, **Qm**, **Rm**, **MPMm**, **Jm**, **Fm**, **Om**.